

MediaCorp NewsBot Election Misinformation

Post-Mortem Report

AI Incident Report Case (C)

Trustworthy AI 51.516 | Week 8 Lab | Team Members: Charles Foo, Sagar Pratap | Date: 16 July 2026

1. Case Summary

NewsBot is a news-summarization LLM, an open-weight foundation model that was fine-tuned in-house by MediaCorp’s 10-person Digital Innovation team. It was trained on a local (Singapore) corpus of news articles, social media posts and chat forums. NewsBot powered the “Trending Now” widget on MediaCorp’s website and app, auto-publishing 60 to 100 short summaries of trending political and general news per day to roughly 2.4 million daily readers during an election season, with no human review before publishing. It generated summaries from parametric memory, with no retrieval grounding and no requirement to cite a source, and its success was measured by summaries published and click-through, not accuracy.

Beginning on Day 4 after launch, NewsBot auto-published fabricated claims about two candidates, asserting that “Candidate A received bribes from foreign governments” and that “Candidate B is not a citizen,” both false. Readers shared the summaries on social media, and they went viral roughly 2.5 weeks before Polling Day, the false claims circulated live for about ten days (Day 4 to Day 14). MediaCorp took the system offline on Day 14, issued a public apology and retractions on Day 16, and an external audit plus an election-commission inquiry followed. The harmed parties were the two candidates (reputational and legal harm), the voting public (exposure to false claims during an election), and MediaCorp itself (trust, legal, and regulatory exposure).

2. Timeline

Dates are anchored based on a fictional election with Polling Day = Nov 4, Report Date = Nov 25

Date	Event
1 Sep 2025	MediaCorp’s internal strategy memo green lights Project NewsBot to reduce editorial workload and improve real-time news summarization on their public website
10 Jan 2026	An approximately 2M document corpus consisting of news articles, social media posts and forum threads was assembled to fine-tune the model on local context. An automated scraping tool was used to grab content that matched trending election-related keywords and phrases, resulting in fringe blogs, a conspiracy forum and three known fake-news domains being included. No source whitelisting or credibility filter was applied.
11 Apr 2026	NeoBot’s pre-launch evaluation checks were conducted using ROUGE-L summarization quality, latency and a toxicity classifier. The evaluation did not include factual accuracy, misinformation read-team or election stress test
20 Sep 2026	Launch was approved against user engagement-lift and toxicity pass criteria. The

	approval chain did not include the editorial leadership
7 Oct 2026	NewsBot goes live on “Trending Now”, auto-publishing 60-100 summaries per day without human oversight
11 Oct 2026	NewsBot publishes the first fabricated summary: Candidate A bribery claim (false)
14 Oct 2026	Second fabricated summary published: Candidate B citizenship claim (false). Readers begin to share on social media
17 Oct 2026	Fabricated claims go viral, an independent fact-checking outlet queries MediaCorp and Candidate A’s campaign sends a legal demand letter
19 Oct 2026	Internal escalation starts within MediaCorp, but the response is ad-hoc due to the absence of a runbook. The ML team confirms the summaries were model-generated, not sourced
21 Oct 2026	NewsBot taken offline, the off-switch requires a product redeploy (not a configuration toggle), adding hours of delay.
23 Oct 2026	MediaCorp issues a public apology and retracts the affected summaries
20 Nov 2026	External audit commissioned, the election commission opens an inquiry into the automated political content

3. Root Cause Analysis

3.1 Technical dimension

Finding: NewsBot summarized from parametric memory with no retrieval grounding and no requirement to cite an actual source, hence it could assert claims with no document backing. Firstly, the dataset used to train NewsBot was not reviewed for quality and legitimacy. Roughly 20% of its 2M document corpus was keyword-scraped content from user forums and aggregator that swept in fringe blogs, conspiracy forums and three known fake-news domains without any credibility filter. Secondly, the model was tuned to optimize human-like summarization (ROUGE-L), fluency reward and fast performance. There were no metrics to measure factuality or claim grounding. Outputs stayed fluent and confident even when fabricated, with no uncertainty signal and no abstain path for low-confidence political claims.

Evidence: Packet Section 3 (Technical lens): corpus mix (40% newswire, 25% archive, 20% social, 15% forums), the unfiltered keyword-scraped slice, the ROUGE-L plus fluency objective with no faithfulness metric, the absence of retrieval and citation, and no calibration or abstain path.

Why the gap existed: Factuality was never an optimization target or a measured metric for NewsBot. The objective rewarded fluent, plausible summaries, while provenance filtering, grounding, hallucination checks and an abstain path when uncertain were all out of scope. To the system, a confident false claim was indistinguishable from a correct one.

3.2 Process dimension

Finding: NewsBot pre-launch testing and evaluation covered only summarization quality, latency and a

toxicity classifier with no factual-accuracy test, misinformation red-team, or election-specific stress test. The go-live was owned by the Head of Digital Product and was approved on the evidence of an A/B user engagement lift plus a toxicity pass. Editorial leadership was excluded from the approval chain. The full launch landed roughly 2.5 weeks before Polling Day with no staged rollout and no soft launch conducted with a focus group to identify real issues that could occur during elections. While NewsBot was live, there was no human-in-the-loop review before news summarizations were published. Additionally, there was also no rollback runbook, so disabling the system required a redeployment rather than a configuration toggle.

Evidence: Packet Section 3(Process lens) and the timeline: Sep 20 go-live record and the Oct 21 redeployment to disable NewsBot entry

Why the gap existed: The go-live optimized for high user engagement and was owned by product. Editorial was out of the approval loop, and no stakeholder was accountable for factual accuracy. “Summarization just rephrases articles” was assumed safe and never tested against hallucination, and no policy required a factuality test, a human gate, or heightened controls during an election. This is a skipped MEASURE function under the NIST AI Risk Management Framework, which calls for testing validity, safety, and accountability, not just quality and tone [2].

3.3 Governance dimension

Finding: No single role owned output quality in production and there was no recognition of an accountable role or team to oversee a publicly safe deployment. The ML team reviewed uptime and latency dashboards but not content quality. There was no content monitoring or alerting on political claims, so the first signal came from external complaints (a fact-checking outlet and a campaign legal letter) rather than internal tooling. No incident-response plan existed for AI-generated content. Comms and legal were not briefed before the rollout, so the incident response was ad-hoc. Summaries were not labelled as AI-generated, as also recommended by the EU AI Act Article 50, which requires AI-generated text informing the public on matters of public interest to be disclosed as such [3], and MediaCorp had no prior engagement with the election commission and no election-integrity policy for automated political content.

Evidence: Packet Section 3 (Governance lens) and the timeline: the Oct 17 external first signal, the Oct 19 ad-hoc escalation, and the November regulator inquiry.

Why the gap existed: No role owned the content safety risk in production, and the system’s KPI was engagement and published volume, not accuracy. Monitoring, disclosure, incident response, and regulator engagement were nobody’s mandate until publicly harm content surfaced externally.

3.4 Primary root cause

Primary root cause: MediaCorp launched a news summarizer that could invent abstract claims, without adequate guardrails for content grounding and safety, and no labelling of AI-generated content. Their

go-live process was not setup to catch the content safety gaps and model hallucination as the focus remained primarily on user engagement uplift and content toxicity pass metrics, with no approval from editorial leadership required and no named owner of production output quality who was accountable for whether the published summaries were factual.

4. Corrective Actions

Dates are anchored based on report issue date 27 Nov 2026

4.1 Short-term (1–2 weeks)

#	Action	Owner	Deadline	Success Metric
1	<u>Keep NewsBot offline and add a single-click toggle kill-switch</u> so the system can be disabled without a production redeployment	Head of Digital Product with IT Application Support/SRE	2 Dec 2026 (1 week)	Verifiable disabling of NewsBot in under 5 mins in a live drill
2	<u>Complete retractions and open regulator cooperation:</u> audit every summary published from Day 0- Day 14 and publish corrections for all falsified claims. Formally respond to both candidates' campaigns and open a channel with the election commission	Editor-in-Chief with Legal and Comms	9 Dec 2026 (2 weeks)	100% of flagged summaries retracted or corrected and documented responses issued to both campaigns and election commission

4.2 Long-term (3–6 months)

#	Action	Owner	Deadline	Success Metric
1	<u>Retrain the model:</u> rebuild the training data corpus using a whitelisted source allow-list and credibility filter reviewed by human. Model fine-tuning should focus on achieving factual grounding. Model summarization should include source-sentence citations. Add an automated claim-verification and fact-check gate to sanitize model output. Include an abstain path for the model to refrain from low-confidence claims, especially political and social in nature.	Chief Data & AI Officer (CDAO)	27 Feb 2027 (3 months)	Pre-launch hallucination rate below a defined threshold on an election test set, with every published summary traceable to cited sources.
2	<u>Close the gaps in the go-live process:</u> mandate a factual accuracy and misinformation red-team pass before any deployment, not just engagement and	Head of Digital Product with Editor-in-Chief	27 Feb 2027 (3 months)	100% of political summaries editor-approved before publish, and no

	toxicity pass. Add human-in-the-loop review so no candidate or election summary publishes without editor sign-off. Stage election rollouts instead of a full-fledged release on Day 1.			political deployment ships without a documented factuality pass
3	<u>Establish decision rights, ownership, monitoring, and election policy</u> : require joint go-live approval from editorial leadership, a named Output Quality Owner and Legal, each with authority to block release. Output Quality owner should be accountable for production content. Run content monitoring with alerting on political claims and introduce an incident response runbook. Label all summaries as AI-generated with a reader report button and suppress auto-generation for high-risk election topics, delegate to human authorship	Editor-in-Chief, VP Product and the named Output Quality Owner	27 May 2027 (6 months)	No deployment ships without documented sign-off from editorial, the Output Quality Owner and Legal. Monitoring is live with appropriate thresholds for escalation and high-risk political topics have been suppressed from model summarization

5. References

- Case C Incident Packet, MediaCorp NewsBot (Trustworthy AI 51.516, Week 8 workshop materials)
- [1] C.-Y. Lin, "ROUGE: A package for automatic evaluation of summaries," in Text Summarization Branches Out, Barcelona, Spain, Jul. 2004, pp. 74–81. [Online]. Available: <https://aclanthology.org/W04-1013.pdf>
- [2] National Institute of Standards and Technology, Artificial Intelligence Risk Management Framework (AI RMF 1.0), NIST AI 100-1, Gaithersburg, MD, USA, Jan. 2023. [Online]. Available: <https://doi.org/10.6028/NIST.AI.100-1>
- [3] European Parliament and Council of the European Union, "Regulation (EU) 2024/1689 of 13 June 2024 laying down harmonized rules on artificial intelligence (Artificial Intelligence Act)," Official Journal of the European Union, L 2024/1689, Jul. 2024. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>
- (Supplementary) Grok AI Election Misinformation- "Secretaries of state urge Musk to fix AI chatbot spreading false election info," The Washington Post, Aug. 4, 2024. [Online]. Available: <https://www.washingtonpost.com/politics/2024/08/04/secretaries-state-urge-musk-fix-ai-chatbot-spreading-false-election-info/>